

Group Name: Three DataMiners Student1: Vinay Huchanahalli Nagaraju(N10180893) Student2:Kalpana Menthem(N10155694)
Student3:Anudeep Gottigundala(N10155678)

In [33]:

```
#Importing required libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
from sklearn.model_selection import GridSearchCV
import pydot
from io import StringIO
from sklearn.tree import export_graphviz
#from dm_tools import data_prep
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, accuracy_score
from collections import Counter
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, accuracy_score
from sklearn.feature_extraction.text import CountVectorizer
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import ClusterCentroids
from PIL import Image
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
import seaborn as sns
from PIL import Image
```

In [4]:

```
# Preprocessing
#listing the nominal and numerical values
nominal_cols = ['Auction', 'Make', 'TopThreeAmericanName', 'Color', 'Transmission', 'Nationality', 'Size', 'VNST', 'WheelType']

num_cols = ['VehYear', 'VehBCost', 'VehOdo', 'IsOnlineSale',
'WarrantyCost', 'MMRAcquisitionAuctionAveragePrice', 'MMRAcquisitionAuctionCleanPrice',
'MMRAcquisitionRetailAveragePrice', 'MMRAcquisitionRetailCleanPrice',
'MMRCurrentAuctionAveragePrice', 'MMRCurrentAuctionCleanPrice',
'MMRCurrentRetailAveragePrice', 'MMRCurrentRetailCleanPrice']
```

In [5]:

```
## Defining missing values
def fill_missing_values(df):

    for col in num_cols:
        df[col] = df[col].fillna(df[col].median())

    for col in nominal_cols:
        mode = df[col].mode()[0]
        df[col] = df[col].fillna(mode)

    return df
```

In [6]:

```
#deleting the unwanted features
def feature_engineering(df):

    del df['WheelTypeID'] #Wheeltype is used. WheelTypeID is dervided from WheelType
    del df['PurchaseID'] # Just a serial number.
    del df['ForSale'] #Just inclined to yes excluding 6 records-No contribution to the prediction
```

```

del df['PurchaseDate'] #We are using Vehicle year to measure the time series and hence,
deleting it.
del df['MMRCurrentRetailRatio'] #Derived from MMRCurrentRetailAveragePrice and
MMRCurrentRetailCleanPrice
del df['PRIMEUNIT'] # more than 80% of the values are undefined(?)
del df['AUCGUART'] # more than 80% of the values are undefined(?)
del df['PurchaseTimestamp'] # We are using year as a way of time measure, derived from
PurchaseDate. We are deleting both of them as we have year to measure.

return df

```

In [7]:

```

def data_type_change(df):

    #Assigning nominal values to either binary or numerical to support the datatype change
    Transmission_map = {"AUTO":0, "MANUAL": 1, 'Manual':1}
    df['Transmission'] = df['Transmission'].map(Transmission_map)
    df['Transmission'].fillna(df['Transmission'].mode(),inplace=True)

    WheelType_map = {"Alloy":1, "Covers": 2, "Special": 3}
    df['WheelType'] = df['WheelType'].map(WheelType_map)

    Auction_map={'ADESA':0, 'MANHEIM':1, 'OTHER':2}
    df['Auction']=df['Auction'].map(Auction_map)

    Make_map={'ACURA':0, 'BUICK':1, 'CADILLAC':3, 'CHEVROLET':4, 'CHRYSLER':5, 'DODGE':6, 'FORD':7, 'GMC':
8, 'HONDA':9, 'HYUNDAI':10, 'INFINITI':11, 'ISUZU':12, 'JEEP':13, 'KIA':14, 'LEXUS':15, 'LINCOLN':16, 'MAZDA
':17, 'MERCURY':18, 'MINI':19, 'MITSUBISHI':20, 'NISSAN':21, 'OLDSMOBILE':22, 'PONTIAC':23, 'SATURN':24, 'S
CION':25, 'SUBARU':26, 'SUZUKI':27, 'TOYOTA':2, 'VOLKSWAGEN':28, 'VOLVO':29}
    df['Make']=df['Make'].map(Make_map)
    #df['Make'].fillna(df['Make'].mode(),inplace=True)

    american_name_map={'CHRYSLER':0, 'FORD':1, 'GM':2, 'OTHER':3}
    df['TopThreeAmericanName']=df['TopThreeAmericanName'].map(american_name_map)

    Color_map={'BEIGE':0, 'BLACK':1, 'BLUE':2, 'BROWN':3, 'GOLD':4, 'GREEN':5, 'GREY':6, 'MAROON':7, 'NOT A
VAIL':8, 'ORANGE':9, 'OTHER':10, 'PURPLE':11, 'RED':12, 'SILVER':13, 'WHITE':14, 'YELLOW':15}
    df['Color']=df['Color'].map(Color_map)

    Nationality_map={'AMERICAN':0, 'OTHER':1, 'OTHER ASIAN':2, 'TOP LINE ASIAN':3, 'USA':4}
    df['Nationality']=df['Nationality'].map(Nationality_map)

    Size_map={'COMPACT':0, 'CROSSOVER':1, 'LARGE':2, 'LARGE SUV':3, 'LARGE TRUCK':4, 'MEDIUM':5, 'MEDIUM
SUV':6, 'SMALL SUV':7, 'SMALL TRUCK':8, 'SPECIALTY':9, 'SPORTS':10, 'VAN':11}
    df['Size']=df['Size'].map(Size_map)

    vnst_map = {'TX':0, 'FL':1, 'CO':2, 'NC':3, 'AZ':4, 'CA':5, 'OK':6, 'SC':7, 'TN':8, 'GA':9, 'VA':10, 'MO
':11, 'PA':12, 'NV':13, 'IN':14, 'MS':15, 'LA':16, 'NJ':17, 'NM':18, 'KY':19, 'AL':20, 'IL':21, 'UT':22, 'WV':2
3, 'WA':24, 'OR':25, 'NH':26, 'NE':27, 'OH':28, 'ID':29, 'NY':30}
    df['VNST'] = df['VNST'].map(vnst_map)

    # changing datatypes as required
    df['Transmission'] = df['Transmission'].astype(float)
    df['Auction'] = df['Auction'].astype(float)
    df['Make'] = df['Make'].astype(float)
    df['TopThreeAmericanName'] = df['TopThreeAmericanName'].astype(float)
    df['Nationality'] = df['Nationality'].astype(float)
    df['Size'] = df['Size'].astype(float)
    df['VNST'] = df['VNST'].astype(float)
    df['VehBCost'] = df['VehBCost'].astype(float)
    df['WheelType'] = df['WheelType'].astype(int)
    df['IsOnlineSale'] = df['IsOnlineSale'].astype(float)
    df['MMRAcquisitionAuctionAveragePrice'] = df['MMRAcquisitionAuctionAveragePrice'].astype(float)
    df['MMRAcquisitionAuctionCleanPrice'] = df['MMRAcquisitionAuctionCleanPrice'].astype(float)
    df['MMRAcquisitionRetailAveragePrice'] = df['MMRAcquisitionRetailAveragePrice'].astype(float)
    df['MMRAcquisitonRetailCleanPrice'] = df['MMRAcquisitonRetailCleanPrice'].astype(float)
    df['MMRCurrentAuctionAveragePrice'] = df['MMRCurrentAuctionAveragePrice'].astype(float)
    df['MMRCurrentAuctionCleanPrice'] = df['MMRCurrentAuctionCleanPrice'].astype(float)
    df['MMRCurrentRetailAveragePrice'] = df['MMRCurrentRetailAveragePrice'].astype(float)
    df['MMRCurrentRetailCleanPrice'] = df['MMRCurrentRetailCleanPrice'].astype(float)

```

```
return df
```

In [8]:

```
# IsOnlineSale is a binary variable which accepts either zero or one. So we are replacing other values with nan, which will be replaced by the median
def error_replacing_For_IsOnlineSale(df):
    mask = df['IsOnlineSale'] == -1
    df.loc[mask, 'IsOnlineSale'] = np.nan
    mask = df['IsOnlineSale'] == 2
    df.loc[mask, 'IsOnlineSale'] = np.nan
    mask = df['IsOnlineSale'] == 4
    df.loc[mask, 'IsOnlineSale'] = np.nan

    return df
```

In [9]:

```
# Converting nominal cols to one-hot vectors
def convert_nominal_cols(df):
    global nominal_cols
    df_with_dummies = pd.get_dummies(df, columns = nominal_cols)

    return df_with_dummies
```

In [10]:

```
def analyse_feature_importance(dm_model, feature_names, n_to_display=20):
    # grab feature importances from the model
    importances = dm_model.feature_importances_

    # sort them out in descending order
    indices = np.argsort(importances)
    indices = np.flip(indices, axis=0)

    # limit to 20 features, you can leave this out to print out everything
    indices = indices[:n_to_display]

    for i in indices:
        print(feature_names[i], ': ', importances[i])

def visualize_decision_tree(dm_model, feature_names, save_name):
    # Visualize the model using three parameters
    import pydot
    from io import StringIO
    from sklearn.tree import export_graphviz

    dotfile = StringIO()
    export_graphviz(dm_model, out_file=dotfile, feature_names=feature_names)
    graph = pydot.graph_from_dot_data(dotfile.getvalue())
    graph[0].write_png(save_name) # saved in the following file
```

In [13]:

```
def preprocessing():
    df = pd.read_csv('CaseStudyData.csv')
    # dropping the columns with continuously 10 null values. We have 44 records with blank values for 26 variables (or columns)
    new1_df = df.dropna(axis=0, thresh=10)
    # Replacing ? with null values
    new_df = new1_df.replace(['?'], np.nan, inplace=False)
    error_replacing_For_IsOnlineSale(new_df)
    fill_missing_values(new_df)
    feature_engineering(new_df)
    convert_nominal_cols(new_df)
    return new_df
```

In [14]:

```
df = pd.read_csv('CaseStudyData.csv')
df = preprocessing()
```

```
df2 = data_type_change(df)

print(df2.info())
Y = df2['IsBadBuy']
X = df2.drop(['IsBadBuy'], axis=1)
rs=10
#Split the data based on training and testing with 70 and 30%
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3, stratify=Y, random_state=r
s)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3020: DtypeWarning:
Columns (27) have mixed types. Specify dtype option on import or set low_memory=False.
interactivity=interactivity, compiler=compiler, result=result)
C:\ProgramData\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3185: DtypeWarning:
Columns (27) have mixed types. Specify dtype option on import or set low_memory=False.
if (yield from self.run_code(code, result)):
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41432 entries, 0 to 41475
Data columns (total 23 columns):
Auction          41432 non-null float64
VehYear          41432 non-null float64
Make             41432 non-null float64
Color            41432 non-null int64
Transmission     41432 non-null float64
WheelType        41432 non-null int32
VehOdo           41432 non-null float64
Nationality      41432 non-null float64
Size             41432 non-null float64
TopThreeAmericanName 41432 non-null float64
MMRAcquisitionAuctionAveragePrice 41432 non-null float64
MMRAcquisitionAuctionCleanPrice 41432 non-null float64
MMRAcquisitionRetailAveragePrice 41432 non-null float64
MMRAcquisitionRetailCleanPrice 41432 non-null float64
MMRCurrentAuctionAveragePrice 41432 non-null float64
MMRCurrentAuctionCleanPrice 41432 non-null float64
MMRCurrentRetailAveragePrice 41432 non-null float64
MMRCurrentRetailCleanPrice 41432 non-null float64
VNST             41432 non-null float64
VehBCost         41432 non-null float64
IsOnlineSale     41432 non-null float64
WarrantyCost     41432 non-null float64
IsBadBuy         41432 non-null int64
dtypes: float64(20), int32(1), int64(2)
memory usage: 7.4 MB
None
```

Task 1. Data Selection and Distribution.

In [15]:

```
print("Proportion of cars who can be classified as a kick : ",df.IsBadBuy.value_counts()[1]/(df.Is
BadBuy.value_counts()[0]+df.IsBadBuy.value_counts()[1]))
```

Proportion of cars who can be classified as a kick : 0.12948928364549142

1. What is the proportion of cars who can be classified as a “kick”? Answer: Nearly 13% of the cars can be classified as Kick by the given data of IsBadBuy. IsBadBuy=1 suggests that it is a Kick.
2. Did you have to fix any data quality problems? Detail them Answer: We have missing values, noise and erroneous values and incorrect format for some of the features in the given data. These are the data quality problems we encountered. Detailed description is given below.

- i) We have 44 records with blank values for 26 features continuously, which of no use to predict a car is a Kick or not..
- ii) We have derived features from other fetures(Eg.WheelTypeID is dervided from WheelTy pe,MMRCurrentRetailRatio derived from MMRCurrentRetailAveragePrice and MMRCurrentRetailCleanPrice and PurchaseTimestamp is derived from PurchaseDate).
- iii) PurchaseID is just a serial number which is not contributing to predict the target .

- iv) PRIMEUNIT and AUCGUART are having undefined values as '?' for more than 80% of the records.
- v) ForSale has just 6 records with value "No", remaining all 'Yes' amongst 41476 records, which is again not useful to predict the target.
- vi) PurchaseDate is nowhere needed as we have VehYear, which is a measure of time to predict the target.
- vii) PurchaseTimestamp is derived from PurchaseDate, and we have VehYear to measure the time. So no need of PurchaseTimestamp to predict a car is Kick or not.
- Viii) Most of the features have skewness as shown below using boxplot and Histograms for all the variables.
- ix) Some of the records of features have undefined value which is a '?' (erroneous values) and also have blank values.

3. Can you identify any clear patterns by initial exploration of the data using histogram or box plot?

We have plotted the histograms using distplot for numerical/categorical values and countplot for nominal values.

Auction "MANHEIM" has more distribution compared to the other auction company names i.e. There are more chances of a kick cars from the MANHEIM company. CHEVROLET cars, GM cars, Silver coloured cars, Auto transmission cars, American nationality cars, Alloy wheel type cars have the highest count, which are kick cars from the histograms. We have skewness for the price variables which means the data is distributed for a particular range of price.

In [16]:



```
# Histogram (Univariate Analysis)
print('Histogram plot for Nominal columns')
for col in nominal_cols:
    dg = sns.countplot(df[col])
    plt.show()
    print('-'*70)



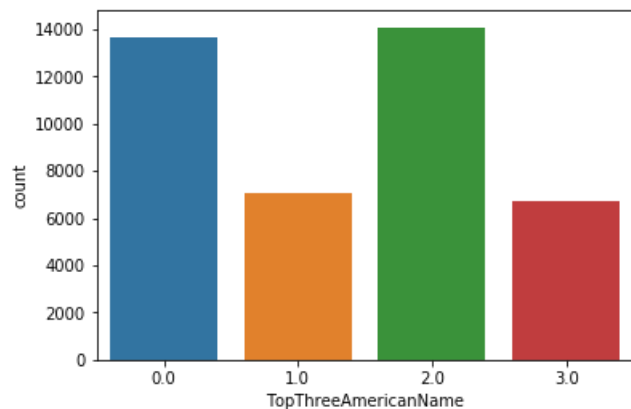

print('='*100)

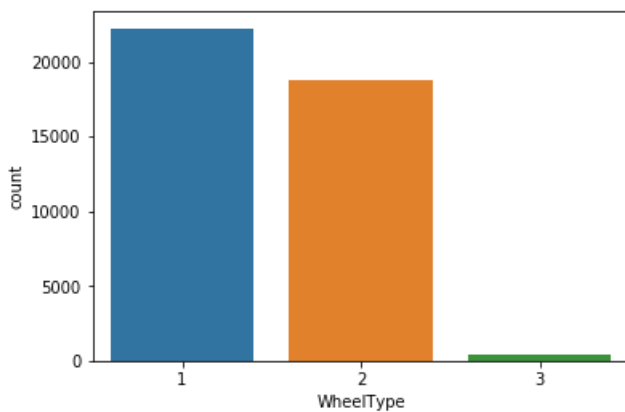



print('Histogram plot for numerical columns')

print('-'*50)
for col in num_cols:
    dg = sns.distplot(df[col])
    plt.show()
    print('-'*70)
```

Histogram plot for Nominal columns



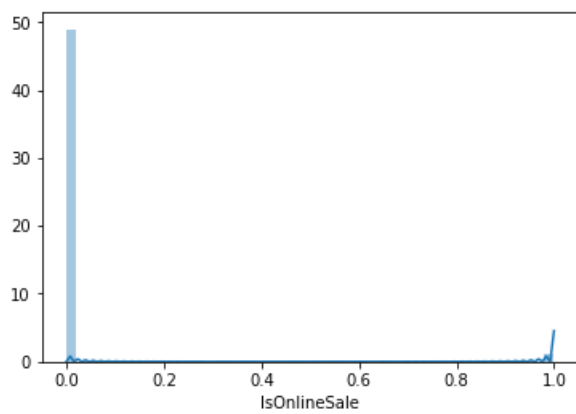


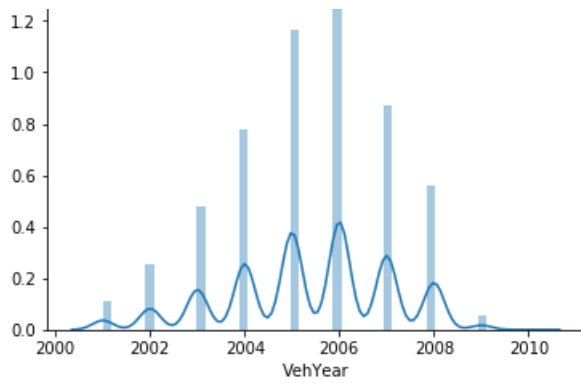






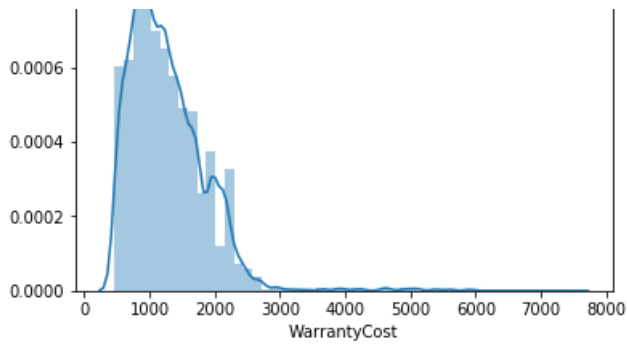







+++++
 +++++
 Histogram plot for numerical columns


```

C:\ProgramData\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
  
```









In [19]:



```
# BoxPlot (Univariate Analysis)


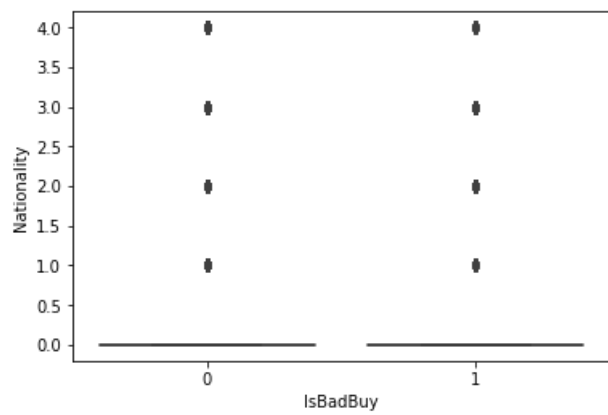
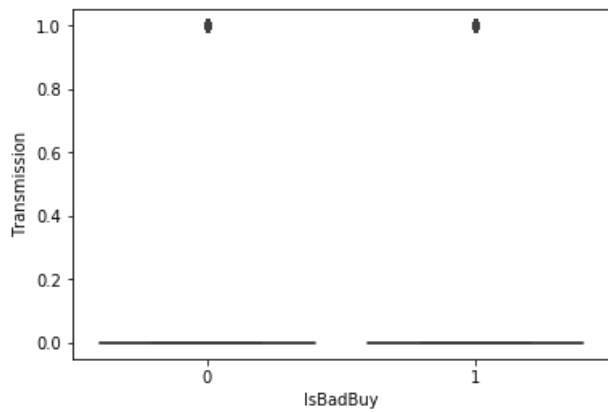


print('Box Plot for nominal columns')
for col in nominal_cols:
    ax = sns.boxplot(x="IsBadBuy", y=col, data=df)
    plt.show()
    print('-----')




print('BoxPlot for numerical columns')

print('-----')
for col in num_cols:
    ax = sns.boxplot(x="IsBadBuy", y=col, data=df)
    print('-'*70)
```


Box Plot for nominal columns







BoxPlot for numerical columns





1. What variables did you include in the analysis and what were their roles and measurement level set? Justify your choice

We are including all the variables except

WheelTypeID,PurchaseID,ForSale,PurchaseDate,MMRCurrentRetailRatio,PRIMEUNIT,AUCGUART and PurchaseTimestamp.

Role: Auction will give the company name which can give kick cars and non-kick cars. From vehicle year, we can predict which year making car is a kick or not. Vehicle make will give the which type of make gives the kick cars and non-kick cars. From Color, we can predict the kick cars based on the color of the car. Transmission gives either Auto cars or manual cars are kicks. WheelType, based on the wheelttype of a car , can distinguish the kick cars from non-kick cars. VehOdo will differntiate the cars based on the kilometers it has run already. Natioanlity also gives the some prediction for the kick cars. Size can also predict the kick cars based on the size of the vehicle.TopThreeAmericanName, can distinguish the kick cars from non-kick cars as it will give the trust. Based on the prices also we can predict the kick cars, as the prices are trustable or not and which price is suitable based on the car. Geograpihc region can affect the kick cars.VehBCost, can predict the kick cars based on the price of the car at the puchase time. IsOnlineSale, says weather the car is available online or not. Warranty cost also affects the kick car.

Measurement Level Set: Below command gives the measurement level for the variables which we included. For example, Transmission=0 means Auto and Transmission=1 means Manual. Auto cars slightly are the kick cars based on the measurement level given below.

1. What distribution scheme did you use? What data partitioning allocation did you set? Explain your selection. Answer: We are dividing the data into training and test sets. First we will train the model using train data and then we will test it using the test data. This is the distribution scheme which has been used here. We have used the 70/30 partition allocation here. 70% of the data has been used for training and 30% has been used for testing. Our data is not a big data so no need of validation data set. That is why we are just using trainig and testing data sets in our distribution scheme. We have taken 70/30 as partitioning allocation as it is a common criteria. We are also using stratify while performing the split to ensure the same ratio of positive and negative targets in both trainig and testing datasets. We can see the command for this below.

In [20]:

```
# Measurement Level Set (Univariate Analysis)

print('Measurement level for nominal columns')
for col in nominal_cols:
    print(df.groupby(['IsBadBuy'])[col].value_counts(normalize=True))
    print('-----')

print('Measurement level for numerical columns')

print('-----')
for col in num_cols:
    print(df.groupby(['IsBadBuy'])[col].value_counts(normalize=True))
    print('-'*70)
```

Measurement level for nominal columns

IsBadBuy	Auction	
0	1.0	0.541825
	0.0	0.258657
	2.0	0.199518
1	1.0	0.489469
	0.0	0.327493
	2.0	0.183038

Name: Auction, dtype: float64

IsBadBuy	Make	
0	4.0	0.238085
	6.0	0.182078
	7.0	0.149361
	5.0	0.126459
	23.0	0.057476
	14.0	0.032329
	24.0	0.029390
	21.0	0.027782
	13.0	0.022874
	10.0	0.022763
	27.0	0.019907
	2.0	0.016414
	30.0	0.013025

	20.0	0.013633
	17.0	0.012505
	18.0	0.011728
	1.0	0.009732
	8.0	0.008595
	9.0	0.006294
	22.0	0.003189
	12.0	0.002079
	25.0	0.001941
	28.0	0.001691
	16.0	0.001026
	11.0	0.000471
	3.0	0.000388
	26.0	0.000388
	0.0	0.000333
	29.0	0.000333
	19.0	0.000305
	15.0	0.000250
1	7.0	0.199627
	4.0	0.179124
	6.0	0.152470
	5.0	0.130103
	23.0	0.052563
	24.0	0.034483
	21.0	0.034296
	14.0	0.031873
	13.0	0.029823
	10.0	0.025349
	27.0	0.023113
	18.0	0.019385
	17.0	0.015098
	2.0	0.013420
	20.0	0.013048
	1.0	0.011556
	8.0	0.007642
	9.0	0.006710
	22.0	0.005778
	16.0	0.003169
	28.0	0.002237
	11.0	0.001864
	19.0	0.001491
	0.0	0.001305
	12.0	0.001305
	25.0	0.001305
	15.0	0.000746
	3.0	0.000559
	26.0	0.000559

Name: Make, dtype: float64

	IsBadBuy	TopThreeAmericanName
0	2.0	0.346938
	0.0	0.331356
	1.0	0.162115
	3.0	0.159592
1	0.0	0.312395
	2.0	0.291705
	1.0	0.222181
	3.0	0.173719

Name: TopThreeAmericanName, dtype: float64

	IsBadBuy	Color
0	13	0.205617
	14	0.166607
	2	0.142318
	1	0.107827
	6	0.103446
	12	0.087864
	4	0.072144
	5	0.042865
	7	0.025148
	0	0.021405
	9	0.006377
	3	0.005822
	11	0.004686
	10	0.003327
	15	0.003299
	8	0.001248
1	12	0.010011

```
1      13      0.210811
      14      0.164212
      2      0.134576
      6      0.096365
      1      0.093756
     12      0.091705
      4      0.085182
      5      0.046598
      7      0.024604
      0      0.022740
      3      0.007269
     11      0.006710
      9      0.004660
     15      0.004101
      8      0.003728
     10      0.002982
```

Name: Color, dtype: float64

```
-----
IsBadBuy  Transmission
0          0.0          0.963512
          1.0          0.036488
1          0.0          0.966449
          1.0          0.033551
```

Name: Transmission, dtype: float64

```
-----
IsBadBuy  Nationality
0          0.0          0.837469
          2.0          0.106857
          3.0          0.050406
          4.0          0.002939
          1.0          0.002329
1          0.0          0.822740
          2.0          0.115564
          3.0          0.054427
          1.0          0.003728
          4.0          0.003541
```

Name: Nationality, dtype: float64

```
-----
IsBadBuy  Size
0          5.0          0.427233
          2.0          0.125405
          6.0          0.106552
          0.0          0.093132
         11.0          0.081044
          4.0          0.046081
          7.0          0.031636
          9.0          0.024870
          1.0          0.023845
          3.0          0.018965
          8.0          0.011617
         10.0          0.009621
1          5.0          0.397763
          6.0          0.135322
          0.0          0.126002
          2.0          0.082945
         11.0          0.082759
          4.0          0.043802
          7.0          0.035601
          3.0          0.027213
          1.0          0.021249
          9.0          0.018826
         10.0          0.014539
          8.0          0.013979
```

Name: Size, dtype: float64

```
-----
IsBadBuy  VNST
0          0.0          0.216819
          1.0          0.127956
          2.0          0.088197
          3.0          0.086644
          4.0          0.080572
          5.0          0.076219
          6.0          0.064713
          7.0          0.041090
          8.0          0.035878
          9.0          0.031441
         10.0          0.026090
         11.0          0.018710
```


	11.0	0.018743
	12.0	0.015887
	13.0	0.012837
	14.0	0.011590
	15.0	0.010259
	16.0	0.008235
	17.0	0.007680
	19.0	0.005961
	18.0	0.005822
	20.0	0.004464
	22.0	0.004270
	21.0	0.004020
	24.0	0.003549
	25.0	0.003438
	23.0	0.003355
	26.0	0.002440
	27.0	0.000693
	28.0	0.000582
	29.0	0.000388
	30.0	0.000166
1	0.0	0.234110
	1.0	0.118360
	5.0	0.096738
	4.0	0.088910
	3.0	0.087418
	2.0	0.082386
	6.0	0.048649
	7.0	0.033551
	8.0	0.032992
	9.0	0.028518
	10.0	0.028332
	12.0	0.023672
	13.0	0.016775
	11.0	0.015284
	14.0	0.012675
	16.0	0.009692
	15.0	0.007829
	17.0	0.007456
	18.0	0.005405
	21.0	0.003728
	20.0	0.003355
	23.0	0.002982
	19.0	0.002796
	25.0	0.002237
	22.0	0.002050
	26.0	0.001678
	24.0	0.001491
	28.0	0.000746
	27.0	0.000186

Name: VNST, dtype: float64

IsBadBuy	WheelType	
0	1	0.512518
	2	0.477001
	3	0.010480
1	1	0.698975
	2	0.290214
	3	0.010811

Name: WheelType, dtype: float64

Measurement level for numerical columns

IsBadBuy	VehYear	
0	2006.0	0.240136
	2005.0	0.207447
	2007.0	0.166662
	2004.0	0.133696
	2008.0	0.108964
	2003.0	0.078049
	2002.0	0.038900
	2001.0	0.015915
	2009.0	0.010203
	2010.0	0.000028
1	2005.0	0.223672
	2004.0	0.180801
	2006.0	0.180615
	2003.0	0.137745
	2007.0	0.000000

2007.0	0.093756
2002.0	0.088723
2008.0	0.046039
2001.0	0.045107
2009.0	0.003541

Name: VehYear, dtype: float64

IsBadBuy	VehBCost
0	7500.0 0.011811
	6500.0 0.007209
	7800.0 0.006654
	7000.0 0.006488
	7200.0 0.006294
	6000.0 0.006183
	8000.0 0.006155
	7100.0 0.005878
	6300.0 0.005795
	7400.0 0.005268
	6400.0 0.005018
	4200.0 0.004963
	6100.0 0.004824
	7700.0 0.004381
	7300.0 0.004298
	8200.0 0.004131
	5500.0 0.004103
	5000.0 0.003882
	6700.0 0.003715
	6600.0 0.003660
	6800.0 0.003660
	6200.0 0.003632
	7600.0 0.003604
	8100.0 0.003577
	5800.0 0.003410
	7900.0 0.003355
	5700.0 0.003327
	6900.0 0.003133
	4175.0 0.003078
	5900.0 0.002828
	...

1	11385.0 0.000186
	11425.0 0.000186
	11430.0 0.000186
	11445.0 0.000186
	11495.0 0.000186
	11500.0 0.000186
	11505.0 0.000186
	11527.0 0.000186
	11595.0 0.000186
	11600.0 0.000186
	11620.0 0.000186
	11645.0 0.000186
	11705.0 0.000186
	11760.0 0.000186
	11785.0 0.000186
	11845.0 0.000186
	11900.0 0.000186
	12025.0 0.000186
	12090.0 0.000186
	12330.0 0.000186
	12590.0 0.000186
	13535.0 0.000186
	18245.0 0.000186
	19000.0 0.000186
	20100.0 0.000186
	28180.0 0.000186
	29795.0 0.000186
	32300.0 0.000186
	38785.0 0.000186
	45469.0 0.000186

Name: VehBCost, Length: 3237, dtype: float64

IsBadBuy	VehOdo
0	50902.0 0.000166
	67756.0 0.000166
	67860.0 0.000166
	71225.0 0.000166
	74671.0 0.000166

76267.0	0.000166
79600.0	0.000166
84675.0	0.000166
59355.0	0.000139
62143.0	0.000139
62277.0	0.000139
63053.0	0.000139
67138.0	0.000139
67622.0	0.000139
67625.0	0.000139
67953.0	0.000139
69089.0	0.000139
69413.0	0.000139
70269.0	0.000139
70334.0	0.000139
71005.0	0.000139
71783.0	0.000139
72101.0	0.000139
73154.0	0.000139
73232.0	0.000139
74538.0	0.000139
74783.0	0.000139
75007.0	0.000139
75064.0	0.000139
75309.0	0.000139

...

1	103531.0	0.000186
	103575.0	0.000186
	103675.0	0.000186
	103806.0	0.000186
	103834.0	0.000186
	103929.0	0.000186
	104125.0	0.000186
	104716.0	0.000186
	104957.0	0.000186
	105313.0	0.000186
	105536.0	0.000186
	105776.0	0.000186
	105989.0	0.000186
	106225.0	0.000186
	106774.0	0.000186
	106885.0	0.000186
	107091.0	0.000186
	107383.0	0.000186
	107741.0	0.000186
	107860.0	0.000186
	108275.0	0.000186
	108486.0	0.000186
	108825.0	0.000186
	109260.0	0.000186
	109348.0	0.000186
	109549.0	0.000186
	109728.0	0.000186
	109848.0	0.000186
	114184.0	0.000186
	115717.0	0.000186

Name: VehOdo, Length: 31051, dtype: float64

	IsBadBuy	IsOnlineSale
0	0.0	0.978041
	1.0	0.021959
1	0.0	0.982293
	1.0	0.017707

Name: IsOnlineSale, dtype: float64

	IsBadBuy	WarrantyCost
0	920.0	0.041257
	1974.0	0.034408
	2152.0	0.030804
	1215.0	0.029279
	1389.0	0.028807
	1155.0	0.025647
	728.0	0.023373
	803.0	0.022514
	1086.0	0.021377
	1503.0	0.020850
	1703.0	0.020323

1243.0	0.020240
569.0	0.019907
1020.0	0.018882
983.0	0.018188
834.0	0.017578
1272.0	0.017301
533.0	0.016913
1623.0	0.016802
754.0	0.015887
853.0	0.015721
1763.0	0.015582
671.0	0.015222
505.0	0.015111
825.0	0.015055
1373.0	0.014917
693.0	0.014279
1506.0	0.014279
975.0	0.014113
1633.0	0.013974

...

1	1137.0	0.000186
	1181.0	0.000186
	1275.0	0.000186
	1301.0	0.000186
	1418.0	0.000186
	1487.0	0.000186
	1557.0	0.000186
	1571.0	0.000186
	1590.0	0.000186
	1610.0	0.000186
	1634.0	0.000186
	1931.0	0.000186
	1944.0	0.000186
	2090.0	0.000186
	2101.0	0.000186
	2141.0	0.000186
	2251.0	0.000186
	2441.0	0.000186
	2499.0	0.000186
	2700.0	0.000186
	2711.0	0.000186
	2799.0	0.000186
	2838.0	0.000186
	2891.0	0.000186
	2976.0	0.000186
	3115.0	0.000186
	3222.0	0.000186
	3298.0	0.000186
	3667.0	0.000186
	6208.0	0.000186

Name: WarrantyCost, Length: 503, dtype: float64

IsBadBuy	MMRAcquisitionAuctionAveragePrice
0	0.0
	0.012144
	5480.0
	0.005545
	6311.0
	0.002551
	5569.0
	0.001941
	7644.0
	0.001941
	7991.0
	0.001941
	7811.0
	0.001858
	4573.0
	0.001636
	7245.0
	0.001636
	6858.0
	0.001469
	8196.0
	0.001469
	6892.0
	0.001414
	7048.0
	0.001331
	7960.0
	0.001331
	7293.0
	0.001275
	5427.0
	0.001248
	7513.0
	0.001248
	3688.0
	0.001192
	6820.0
	0.001164
	7541.0
	0.001137
	7314.0
	0.001109
	7533.0
	0.001109
	8268.0
	0.001109
	8194.0
	0.001081

	6733.0	0.001054
	5500.0	0.001026
	6867.0	0.000998
	8012.0	0.000998
	6948.0	0.000943
	7171.0	0.000943
	...	
1	14804.0	0.000186
	15091.0	0.000186
	15313.0	0.000186
	15314.0	0.000186
	15585.0	0.000186
	15613.0	0.000186
	15852.0	0.000186
	16412.0	0.000186
	16536.0	0.000186
	16982.0	0.000186
	17229.0	0.000186
	17246.0	0.000186
	18181.0	0.000186
	18843.0	0.000186
	18900.0	0.000186
	19190.0	0.000186
	19250.0	0.000186
	19480.0	0.000186
	19546.0	0.000186
	19810.0	0.000186
	20635.0	0.000186
	21611.0	0.000186
	21870.0	0.000186
	23031.0	0.000186
	25033.0	0.000186
	27680.0	0.000186
	28077.0	0.000186
	28354.0	0.000186
	32250.0	0.000186
	33543.0	0.000186
Name: MMRAcquisitionAuctionAveragePrice, Length: 12518, dtype: float64		

IsBadBuy	MMRAcquisitionAuctionCleanPrice	
0	0.0	0.009954
	6461.0	0.005601
	7450.0	0.002634
	1.0	0.002190
	8892.0	0.001969
	8258.0	0.001913
	6584.0	0.001830
	8449.0	0.001664
	7837.0	0.001553
	9044.0	0.001553
	8107.0	0.001469
	8469.0	0.001442
	5967.0	0.001414
	8466.0	0.001331
	7614.0	0.001248
	8151.0	0.001164
	6235.0	0.001137
	6508.0	0.001137
	9045.0	0.001137
	4783.0	0.001109
	7934.0	0.001109
	7195.0	0.001081
	8006.0	0.001081
	9772.0	0.001081
	7771.0	0.001054
	8187.0	0.001054
	8287.0	0.001054
	9027.0	0.001054
	6920.0	0.001026
	7280.0	0.001026
	...	
1	17369.0	0.000186
	17383.0	0.000186
	17530.0	0.000186
	17625.0	0.000186
	17699.0	0.000186
	17734.0	0.000186

18034.0	0.000186
18305.0	0.000186
18384.0	0.000186
18427.0	0.000186
19745.0	0.000186
20042.0	0.000186
20795.0	0.000186
20809.0	0.000186
21049.0	0.000186
21221.0	0.000186
21234.0	0.000186
21338.0	0.000186
21597.0	0.000186
21605.0	0.000186
23021.0	0.000186
23751.0	0.000186
23969.0	0.000186
25681.0	0.000186
28053.0	0.000186
29498.0	0.000186
30114.0	0.000186
30408.0	0.000186
35215.0	0.000186
36701.0	0.000186

Name: MMRAcquisitionAuctionCleanPrice, Length: 13313, dtype: float64

IsBadBuy	MMRAcquisitionRetailAveragePrice	
0	0.0	0.012144
	6418.0	0.005573
	7316.0	0.002551
	8756.0	0.002052
	11114.0	0.001885
	6515.0	0.001858
	11882.0	0.001580
	7907.0	0.001525
	8325.0	0.001497
	5439.0	0.001442
	9352.0	0.001442
	9097.0	0.001331
	11006.0	0.001303
	7943.0	0.001220
	6361.0	0.001192
	4483.0	0.001164
	8600.0	0.001164
	8614.0	0.001164
	9429.0	0.001137
	7866.0	0.001081
	10574.0	0.001081
	7916.0	0.001054
	9350.0	0.001054
	6378.0	0.001026
	7772.0	0.001026
	10856.0	0.001026
	10875.0	0.001026
	11396.0	0.001026
	9091.0	0.000998
	6440.0	0.000970
		...
1	18619.0	0.000186
	18841.0	0.000186
	18940.0	0.000186
	19107.0	0.000186
	19126.0	0.000186
	19159.0	0.000186
	20008.0	0.000186
	20135.0	0.000186
	20736.0	0.000186
	20737.0	0.000186
	20850.0	0.000186
	20912.0	0.000186
	20976.0	0.000186
	21225.0	0.000186
	21290.0	0.000186
	21336.0	0.000186
	21538.0	0.000186
	21895.0	0.000186
	22786.0	0.000186

23361.0	0.000186
23456.0	0.000186
23840.0	0.000186
24120.0	0.000186
27295.0	0.000186
30048.0	0.000186
30196.0	0.000186
31599.0	0.000186
33872.0	0.000186
35330.0	0.000186
36726.0	0.000186

Name: MMRAcquisitionRetailAveragePrice, Length: 14426, dtype: float64

IsBadBuy	MMRAcquisitonRetailCleanPrice	
0	0.0	0.012116
	7478.0	0.005601
	9722.0	0.002717
	8546.0	0.002495
	11562.0	0.001885
	10103.0	0.001858
	7611.0	0.001830
	9643.0	0.001525
	10268.0	0.001525
	8964.0	0.001442
	9256.0	0.001442
	6944.0	0.001414
	12239.0	0.001414
	11599.0	0.001220
	11513.0	0.001192
	8271.0	0.001164
	11443.0	0.001164
	11447.0	0.001164
	7529.0	0.001137
	5666.0	0.001081
	9303.0	0.001081
	11054.0	0.001081
	10269.0	0.001054
	12565.0	0.001054
	9613.0	0.001026
	10748.0	0.001026
	7234.0	0.000998
	7974.0	0.000998
	9342.0	0.000998
	9624.0	0.000998

1	20732.0	...
	20868.0	0.000186
	21171.0	0.000186
	21370.0	0.000186
	21662.0	0.000186
	21825.0	0.000186
	22145.0	0.000186
	22888.0	0.000186
	23011.0	0.000186
	23123.0	0.000186
	23233.0	0.000186
	23419.0	0.000186
	23433.0	0.000186
	23545.0	0.000186
	23738.0	0.000186
	23825.0	0.000186
	23833.0	0.000186
	24870.0	0.000186
	25363.0	0.000186
	25640.0	0.000186
	25799.0	0.000186
	26151.0	0.000186
	26387.0	0.000186
	29981.0	0.000186
	32383.0	0.000186
	32760.0	0.000186
	33736.0	0.000186
	36096.0	0.000186
	38532.0	0.000186
	40137.0	0.000186

Name: MMRAcquisitonRetailCleanPrice, Length: 14948, dtype: float64

IsBadBuy	MMRCurrentAuctionAveragePrice
0	0.0
	6074.0
	5480.0
	6311.0
	7269.0
	8186.0
	8033.0
	7644.0
	5569.0
	6858.0
	6966.0
	8196.0
	6814.0
	7524.0
	6967.0
	7608.0
	7612.0
	7901.0
	5033.0
	7495.0
	8568.0
	6892.0
	8018.0
	8140.0
	8268.0
	4573.0
	7457.0
	7661.0
	5760.0
	7927.0

	...
1	14827.0
	14982.0
	15038.0
	15161.0
	15292.0
	15368.0
	15581.0
	15605.0
	15852.0
	16091.0
	16412.0
	16645.0
	16721.0
	16988.0
	17240.0
	17343.0
	17779.0
	17844.0
	18416.0
	18546.0
	19359.0
	19963.0
	20817.0
	21940.0
	23015.0
	27543.0
	27795.0
	28099.0
	32250.0
	33369.0

Name: MMRCurrentAuctionAveragePrice, Length: 12440, dtype: float64

IsBadBuy	MMRCurrentAuctionCleanPrice
0	7324.0
	0.0
	6461.0
	7450.0
	1.0
	8892.0
	7898.0
	8107.0
	6584.0
	9279.0
	9044.0
	8484.0

	7500.0	0.001164
	7560.0	0.001137
	8282.0	0.001137
	9237.0	0.001137
	5967.0	0.001109
	8277.0	0.001054
	9325.0	0.001054
	8438.0	0.000998
	8513.0	0.000998
	8811.0	0.000970
	9129.0	0.000970
	8639.0	0.000943
	8669.0	0.000943
	9209.0	0.000943
	7783.0	0.000915
	7885.0	0.000915
	8132.0	0.000915
	8168.0	0.000915
	...	
1	17267.0	0.000186
	17366.0	0.000186
	17432.0	0.000186
	17515.0	0.000186
	17621.0	0.000186
	17625.0	0.000186
	17699.0	0.000186
	17751.0	0.000186
	17767.0	0.000186
	18427.0	0.000186
	18762.0	0.000186
	18942.0	0.000186
	19025.0	0.000186
	19080.0	0.000186
	19760.0	0.000186
	19769.0	0.000186
	20133.0	0.000186
	20422.0	0.000186
	20790.0	0.000186
	20881.0	0.000186
	21514.0	0.000186
	21601.0	0.000186
	21870.0	0.000186
	24293.0	0.000186
	26168.0	0.000186
	29042.0	0.000186
	29811.0	0.000186
	30136.0	0.000186
	35215.0	0.000186
	36478.0	0.000186

Name: MMRCurrentAuctionCleanPrice, Length: 13194, dtype: float64

IsBadBuy	MMRCurrentRetailAveragePrice	
0	0.0	0.006959
	8704.0	0.005213
	6418.0	0.004298
	7316.0	0.002052
	8756.0	0.001747
	10834.0	0.001580
	11674.0	0.001525
	6515.0	0.001442
	7907.0	0.001331
	9352.0	0.001303
	11237.0	0.001164
	10921.0	0.001137
	9753.0	0.001081
	10564.0	0.000970
	11640.0	0.000970
	11710.0	0.000970
	5439.0	0.000943
	7943.0	0.000943
	9429.0	0.000943
	10481.0	0.000887
	11598.0	0.000887
	8885.0	0.000860
	11913.0	0.000860
	9332.0	0.000832
	11713.0	0.000832

	6721.0	0.000804
	8554.0	0.000804
	9005.0	0.000804
	9085.0	0.000804
	9128.0	0.000804
	...	
1	18073.0	0.000186
	18225.0	0.000186
	18559.0	0.000186
	19092.0	0.000186
	19127.0	0.000186
	19160.0	0.000186
	19299.0	0.000186
	19643.0	0.000186
	20090.0	0.000186
	20291.0	0.000186
	20379.0	0.000186
	20912.0	0.000186
	20930.0	0.000186
	21181.0	0.000186
	21431.0	0.000186
	21688.0	0.000186
	22417.0	0.000186
	22581.0	0.000186
	22850.0	0.000186
	23327.0	0.000186
	24286.0	0.000186
	24349.0	0.000186
	24501.0	0.000186
	27269.0	0.000186
	28050.0	0.000186
	29921.0	0.000186
	31128.0	0.000186
	32928.0	0.000186
	35330.0	0.000186
	36539.0	0.000186

Name: MMRCurrentRetailAveragePrice, Length: 14229, dtype: float64

IsBadBuy	MMRCurrentRetailCleanPrice	
0	0.0	0.006959
	10090.0	0.005601
	7478.0	0.004381
	8546.0	0.002024
	10103.0	0.001747
	12387.0	0.001580
	7611.0	0.001497
	11413.0	0.001497
	12864.0	0.001331
	10268.0	0.001303
	9256.0	0.001275
	11706.0	0.001275
	11739.0	0.001109
	11542.0	0.001081
	12701.0	0.001081
	10571.0	0.001054
	12308.0	0.001054
	6944.0	0.001026
	9830.0	0.001026
	11431.0	0.001026
	11944.0	0.000998
	12309.0	0.000998
	11054.0	0.000970
	12272.0	0.000943
	12687.0	0.000943
	9613.0	0.000887
	10599.0	0.000887
	12252.0	0.000860
	7826.0	0.000832
	7234.0	0.000804

	...	
1	20673.0	0.000186
	20806.0	0.000186
	21047.0	0.000186
	21102.0	0.000186
	21370.0	0.000186
	21910.0	0.000186
	22053.0	0.000186

22079.0	0.000186
22120.0	0.000186
22202.0	0.000186
22888.0	0.000186
23285.0	0.000186
23673.0	0.000186
23922.0	0.000186
24113.0	0.000186
24396.0	0.000186
24653.0	0.000186
25060.0	0.000186
25342.0	0.000186
25518.0	0.000186
25970.0	0.000186
26143.0	0.000186
26164.0	0.000186
30194.0	0.000186
31317.0	0.000186
31744.0	0.000186
33014.0	0.000186
35366.0	0.000186
38532.0	0.000186
39896.0	0.000186

Name: MMRCurrentRetailCleanPrice, Length: 14699, dtype: float64

Decision tree

Task 2. Predictive Modeling Using Decision Trees

1. Python: Build a decision tree using the default setting.

Answer the followings: a. What is the classification accuracy on training and test datasets?

As the data we have is imbalanced (we don't have equal proportion of zeros and ones for target variable IsBadBuy), we are doing oversampling and undersampling to balance the dataset. Below are the test and training accuracies for normal data given, oversampled data and undersampled data respectively in the tabular format.

Normal dataset	Oversampling dataset	Under Sampling Dataset
----------------	----------------------	------------------------

Train Accuracy 1.0	1.0	1.0	Test Accuracy 0.7851166532582462	0.7839098954143202	0.22751407884151248
--------------------	-----	-----	----------------------------------	--------------------	---------------------

- b. What is the size of tree (i.e. number of nodes)?

Normal dataset	Oversampling dataset	Under Sampling Dataset
----------------	----------------------	------------------------

NumberOfNodes 7373	7973	923
--------------------	------	-----

- c. How many leaves are in the tree that is selected based on the validation data set?

Normal dataset	Oversampling dataset	Under Sampling Dataset
----------------	----------------------	------------------------

NumberOfNodes 3687	3987	462
--------------------	------	-----

- d. Which variable is used for the first split? What are the competing splits for this first split? WheelType is the variable used for the first split. Auction and VehYear are the competing splits.

- e. What are the 5 important variables in building the tree?

WheelType, Auction, VehYear, Make, TopThreeAmericanName are the five important features here. We can see this from feature importance cell.

- f. Report if you see any evidence of model overfitting.

From the graph, we can observe the test data is performing better than training data from maxdepth 2.5 to 8. After maxdepth=8, test data is not performing better than the training data. So overfitting is there.

- g. Did changing the default setting (i.e., only focus on changing the setting of the number of splits to create a node) help improving

the model? Answer the above questions on the best performing tree.

Answer: Yes. We changed the number of splits by changing the max-depth to 3. Then accuracy has been improved a little bit.

	Normal dataset	Oversampling dataset	Under Sampling Dataset	Max depth i
s 3				

Train Accuracy 1.0 1.0 1.0 0.7325820889 Test Accuracy 0.7851166532582462 0.7839098954143202 0.22751407884151248 0.7746580852

In [21]:

```
# simple decision tree training
model = DecisionTreeClassifier(random_state=rs)
model.fit(X_train, Y_train)
#print Y_train.value_counts()
print ("*****Simple decisoin tree*****")
print("Train accuracy:", model.score(X_train, Y_train))
print("Test accuracy:", model.score(X_test, Y_test))
print("Number of nodes: ",model.tree_.node_count)
Y_pred = model.predict(X_test)
print(classification_report(Y_test, Y_pred))

def get_num_leaves(model):
    n_nodes = model.tree_.node_count
    ll = model.tree_.children_left
    rl = model.tree_.children_right
    count = 0
    for i in range(0,n_nodes):
        if (ll[i] & rl[i]) == -1:
            count = count + 1
    return count

print("Number of leaves in the tree",get_num_leaves(model))

print("Number of leaves present in the leftside",model.tree_.children_left)
print("Number of leaves present in the rightside",model.tree_.children_right)
```

*****Simple decisoin tree*****

Train accuracy: 1.0

Test accuracy: 0.7851166532582462

Number of nodes: 7373

	precision	recall	f1-score	support
0	0.88	0.87	0.88	10820
1	0.21	0.23	0.22	1610
micro avg	0.79	0.79	0.79	12430
macro avg	0.55	0.55	0.55	12430
weighted avg	0.80	0.79	0.79	12430

Number of leaves in the tree 3687

Number of leaves present in the leftside [1 2 3 ... 7371 -1 -1]

Number of leaves present in the rightside [4520 555 138 ... 7372 -1 -1]

In [22]:

```
# simple decision tree training for Oversampled data

sm = SMOTE(random_state=42)
x_res, y_res = sm.fit_sample(X_train, Y_train)
model.fit(x_res, y_res)
print ("*****Prediction for test data*****")
print("Train accuracy:", model.score(x_res, y_res))
print("Test accuracy:", model.score(X_test, Y_test))
print("Number of nodes: ",model.tree_.node_count)
y_pred = model.predict(X_test)
print(classification_report(Y_test, y_pred))

def get_num_leaves(model):
    n_nodes = model.tree_.node_count
    ll = model.tree_.children_left
    rl = model.tree_.children_right
    count = 0
```

```

    for i in range(0,n_nodes):
        if (ll[i] & rl[i]) == -1:
            count = count + 1
    return count

print("Number of leaves in the tree",get_num_leaves(model))

print("Number of leaves present in the leftside",model.tree_.children_left)
print("Number of leaves present in the rightside",model.tree_.children_right)

```

*****Prediction for test data*****

Train accuracy: 1.0

Test accuracy: 0.7839098954143202

Number of nodes: 7973

	precision	recall	f1-score	support
0	0.89	0.86	0.87	10820
1	0.22	0.26	0.24	1610
micro avg	0.78	0.78	0.78	12430
macro avg	0.55	0.56	0.56	12430
weighted avg	0.80	0.78	0.79	12430

Number of leaves in the tree 3987

Number of leaves present in the leftside [1 2 3 ... -1 -1 -1]

Number of leaves present in the rightside [5028 5027 1690 ... -1 -1 -1]

In [23]:

```

# simple decision tree training for undersampled data
cc = ClusterCentroids(random_state=0)
X_under, Y_under = cc.fit_resample(X_train, Y_train)
model = DecisionTreeClassifier(random_state=rs)
model.fit(X_under, Y_under)
print ("*****Decisoin tree with underfitting of the train
data*****")
print("Train accuracy:", model.score(X_under, Y_under))
print("Test accuracy:", model.score(X_test, Y_test))
print("Number of nodes: ",model.tree_.node_count)
print ("*****Prediction for test data*****")
y_pred = model.predict(X_test)
print(classification_report(Y_test, y_pred))

def get_num_leaves(model):
    n_nodes = model.tree_.node_count
    ll = model.tree_.children_left
    rl = model.tree_.children_right
    count = 0
    for i in range(0,n_nodes):
        if (ll[i] & rl[i]) == -1:
            count = count + 1
    return count

print("Number of leaves in the tree",get_num_leaves(model))

```

*****Decisoin tree with underfitting of the train data*****

Train accuracy: 1.0

Test accuracy: 0.22751407884151248

Number of nodes: 923

*****Prediction for test data*****

	precision	recall	f1-score	support
0	0.89	0.13	0.22	10820
1	0.13	0.90	0.23	1610
micro avg	0.23	0.23	0.23	12430
macro avg	0.51	0.51	0.23	12430
weighted avg	0.79	0.23	0.22	12430

Number of leaves in the tree 462

In [24]:

```

# grab feature importances from the model and feature name from the original X

```

```

importances = model.feature_importances_
feature_names = X.columns

# sort them out in descending order
indices = np.argsort(importances)
indices = np.flip(indices, axis=0)

# limit to 20 features, you can leave this out to print out everything
indices = indices[:20]

for i in indices:
    print(feature_names[i], ': ', importances[i])

```

```

WheelType : 0.5193201112555421
MMRCurrentRetailCleanPrice : 0.08695262600861428
Color : 0.05021582558505198
VehOdo : 0.040198242192698845
VehBCost : 0.03152407246499129
WarrantyCost : 0.030431006723877545
MMRAcquisitionAuctionCleanPrice : 0.025506270772206595
MMRCurrentAuctionCleanPrice : 0.024304890248503425
MMRAcquisitionRetailAveragePrice : 0.023707795484547933
Auction : 0.022296079407274006
Size : 0.020683431645761308
MMRCurrentRetailAveragePrice : 0.018971382791140013
MMRAcquisitionRetailCleanPrice : 0.01883585753469289
MMRAcquisitionAuctionAveragePrice : 0.015889166999467193
VNST : 0.014541274111661815
VehYear : 0.013764994799570036
MMRCurrentAuctionAveragePrice : 0.013409476585198204
TopThreeAmericanName : 0.011878920097363115
Make : 0.010779762872448508
Nationality : 0.0025594011596332213

```

In [25]:

```

#retrain with a small max_depth limit

model = DecisionTreeClassifier(max_depth=3, random_state=rs)
model.fit(x_res, y_res)

print("Train accuracy:", model.score(x_res, y_res))
print("Test accuracy:", model.score(X_test, Y_test))

y_pred = model.predict(X_test)
print(classification_report(Y_test, y_pred))

```

```

Train accuracy: 0.7325820889610647
Test accuracy: 0.7746580852775543

```

	precision	recall	f1-score	support
0	0.88	0.85	0.87	10820
1	0.20	0.25	0.22	1610
micro avg	0.77	0.77	0.77	12430
macro avg	0.54	0.55	0.54	12430
weighted avg	0.80	0.77	0.78	12430

In [26]:

```

# grab feature importance from the model and feature name from the original X
importances = model.feature_importances_
feature_names = X.columns

# sort them out in descending order
indices = np.argsort(importances)
indices = np.flip(indices, axis=0)

# limit to 20 features, you can leave this out to print out everything
indices = indices[:20]

for i in indices:
    print(feature_names[i], ': ', importances[i])

```

```

print(feature_names[1], ': ', importances[1])

# visualize
print("Number of nodes: ", model.tree_.node_count)
dotfile = StringIO()
export_graphviz(model, out_file=dotfile, feature_names=X.columns)
graph = pydot.graph_from_dot_data(dotfile.getvalue())
graph[0].write_png("week3_dt_viz.png") # saved in the following file

```

```

WheelType : 0.8196587486514747
Auction : 0.0877895719305143
VehYear : 0.08725508168771123
Nationality : 0.005296597730299794
IsOnlineSale : 0.0
Make : 0.0
Color : 0.0
Transmission : 0.0
VehOdo : 0.0
Size : 0.0
TopThreeAmericanName : 0.0
WarrantyCost : 0.0
MMRAcquisitionAuctionCleanPrice : 0.0
MMRAcquisitionRetailAveragePrice : 0.0
MMRAcquisitionRetailCleanPrice : 0.0
MMRCurrentAuctionAveragePrice : 0.0
MMRCurrentAuctionCleanPrice : 0.0
MMRCurrentRetailAveragePrice : 0.0
MMRCurrentRetailCleanPrice : 0.0
VNST : 0.0
Number of nodes: 13

```

In [27]:

```

test_score = []
train_score = []

# check the model performance for max depth from 2-20
for max_depth in range(2, 21):
    model = DecisionTreeClassifier(max_depth=max_depth, random_state=rs)
    model.fit(x_res, y_res)
    test_score.append(model.score(X_test, Y_test))
    train_score.append(model.score(x_res, y_res))

```


In [28]:

```

# plot max depth hyperparameter values vs training and test accuracy score
print("Number of nodes: ", model.tree_.node_count)
plt.plot(range(2, 21), train_score, 'b', range(2, 21), test_score, 'r')
plt.xlabel('max_depth\nBlue = training acc. Red = test acc.')
plt.ylabel('accuracy')
plt.show()

```

Number of nodes: 6075



In [37]:

```
# visualize
dotfile = StringIO()

analyse_feature_importance(model, X.columns, 20)
visualize_decision_tree(model, X.columns, "optimal_tree.png")
img = Image.open('optimal_tree.png')
new_width = 70000
new_height = 6000
img = img.resize((new_width, new_height), Image.ANTIALIAS)
img.save('optimal_tree.png')
img.show()
```

```
WheelType : 0.27254852450362715
Auction : 0.1799555474226916
VehYear : 0.1292813916062509
Make : 0.06268755690544962
TopThreeAmericanName : 0.05327562753257296
VehOdo : 0.03808119317710718
VehBCost : 0.03538142037570167
MMRCurrentAuctionCleanPrice : 0.022310905812623107
MMRCurrentRetailAveragePrice : 0.021712119423501562
MMRAcquisitionRetailCleanPrice : 0.020259026348441542
MMRCurrentRetailCleanPrice : 0.019790010254355634
MMRAcquisitionRetailAveragePrice : 0.01965782567957907
WarrantyCost : 0.01902925885020823
MMRCurrentAuctionAveragePrice : 0.018443183697614184
MMRAcquisitionAuctionAveragePrice : 0.01806431330321963
MMRAcquisitionAuctionCleanPrice : 0.016178665179954733
Color : 0.016035190805012162
VNST : 0.015816419285077397
Size : 0.011886114811957038
Nationality : 0.006236801846077701
```

Finding optimal hyperparameters with GridSearchCV

1. Python: Build another decision tree tuned with GridSearchCV

a. What is the classification accuracy on training and test datasets? We are considering the oversampled data as it is balancing the target values. Below is the accuracy for training data and testing data for oversampled data with gridsearchCV.

Train accuracy: 0.8312670812373747 Test accuracy: 0.8662107803700724

b. What is the size of tree (i.e. number of nodes)? Is the size different from the maximal tree or the tree in the previous step? Why?

We have number of leaves 707. The size is different when compared to the maximal tree in the previous step because we are using gridsearchCV with different parameters like criterion, max_depth and min_samples_leaf. It will get optimal tree by the models from n possible cases of training set and test sets.

c. How many leaves are in the tree that is selected based on the validation dataset?

{'criterion': 'gini', 'max_depth': 11, 'min_samples_leaf': 10} Number of leaves in the tree 354

d. Which variable is used for the first split? What are the competing splits for this first split?

WheelType has been used for the first split. Competing split is based on the Auction.

e. What are the 5 important variables in building the tree? Five important features are

WheelType Auction VehYear Make TopThreeAmericanName

f. Report if you see any evidence of model overfitting. There is no sign of Overfitting. Because the training and testing accuracy is performing well. The Depth of the tree and number of leaves is optimal. Testing accuracy is also fine.

g. What are the parameters used? Explain your choices. Criterion, max_depth and min_samples_leaf are the parameters used. We are trying to get the optimal tree by choosing the best criterion among gini and entropy and proper max_depth and minimum samples leaf. Criterion is the function to measure the quality of a split.

In [41]:

```
# grid search CV
params = {'criterion': ['gini', 'entropy'],
```



```

        'max_depth': range(2, 12),
        'min_samples_leaf': range(10, 40, 5)}

cv = GridSearchCV(param_grid=params, estimator=DecisionTreeClassifier(random_state=rs), cv=10)
cv.fit(x_res, y_res)

print("Train accuracy:", cv.score(x_res, y_res))
print("Test accuracy:", cv.score(X_test, Y_test))

# test the best model
y_pred = cv.predict(X_test)
print(classification_report(Y_test, Y_pred))

# print parameters of the best model
print(cv.best_params_)

def get_num_leaves(model):
    n_nodes = model.tree_.node_count
    ll = model.tree_.children_left
    rl = model.tree_.children_right
    count = 0
    for i in range(0, n_nodes):
        if (ll[i] & rl[i]) == -1:
            count = count + 1
    return count

print("Number of nodes: ", cv.best_estimator_.tree_.node_count)
print('-'*50)
print("Number of leaves in the tree", get_num_leaves(cv.best_estimator_))
print('-'*50)

```

Train accuracy: 0.8881253218204143

Test accuracy: 0.8462590506838295

	precision	recall	f1-score	support
0	0.88	0.87	0.88	10820
1	0.21	0.23	0.22	1610
micro avg	0.79	0.79	0.79	12430
macro avg	0.55	0.55	0.55	12430
weighted avg	0.80	0.79	0.79	12430

```
{'criterion': 'gini', 'max_depth': 11, 'min_samples_leaf': 10}
```

Number of nodes: 707

Number of leaves in the tree 354

In [42]:

```

# grid search CV
params = {'criterion': ['gini', 'entropy'],
        'max_depth': range(10, 12),
        'min_samples_leaf': range(8, 15, 2)}

cv = GridSearchCV(param_grid=params, estimator=DecisionTreeClassifier(random_state=rs), cv=10)
cv.fit(x_res, y_res)

print("Train accuracy:", cv.score(x_res, y_res))
print("Test accuracy:", cv.score(X_test, Y_test))

# test the best model
y_pred = cv.predict(X_test)
print(classification_report(Y_test, Y_pred))

# print parameters of the best model
print(cv.best_params_)

print("Number of nodes: ", cv.best_estimator_.tree_.node_count)
print('-'*50)
print("Number of leaves in the tree", get_num_leaves(cv.best_estimator_))
print('-'*50)

```

Train accuracy: 0.8881253218204143

```

train accuracy: 0.8881233218204143
Test accuracy: 0.8462590506838295
      precision    recall  f1-score   support

    0         0.88      0.87      0.88     10820
    1         0.21      0.23      0.22      1610

 micro avg       0.79      0.79      0.79     12430
 macro avg       0.55      0.55      0.55     12430
weighted avg       0.80      0.79      0.79     12430

{'criterion': 'gini', 'max_depth': 11, 'min_samples_leaf': 10}
Number of nodes: 707
-----
Number of leaves in the tree 354
-----

```

In [43]:

```

import numpy as np
import pydot
from io import StringIO
from sklearn.tree import export_graphviz

def analyse_feature_importance(dm_model, feature_names, n_to_display=20):
    # grab feature importances from the model
    importances = dm_model.feature_importances_

    # sort them out in descending order
    indices = np.argsort(importances)
    indices = np.flip(indices, axis=0)

    # limit to 20 features, you can leave this out to print out everything
    indices = indices[:n_to_display]

    for i in indices:
        print(feature_names[i], ': ', importances[i])

def visualize_decision_tree(dm_model, feature_names, save_name):
    dotfile = StringIO()
    export_graphviz(dm_model, out_file=dotfile, feature_names=feature_names)
    graph = pydot.graph_from_dot_data(dotfile.getvalue())
    graph[0].write_png(save_name) # saved in the following file

analyse_feature_importance(cv.best_estimator_, X.columns, 20)
visualize_decision_tree(cv.best_estimator_, X.columns, "optimal_tree.png")
img = Image.open('optimal_tree.png')
new_width = 70000
new_height = 6000
img = img.resize((new_width, new_height), Image.ANTIALIAS)
img.save('optimal_tree.png')
img.show()

```

```

WheelType : 0.37947729011943926
Auction : 0.24738203591539104
VehYear : 0.15653290048457935
Make : 0.06890074370292999
TopThreeAmericanName : 0.05672496040350978
VehBCost : 0.01503433953486562
VehOdo : 0.011850695926835592
MMRAcquisitionRetailAveragePrice : 0.007114144425115366
MMRCurrentAuctionCleanPrice : 0.006725079060459914
MMRCurrentRetailCleanPrice : 0.006560750300985196
MMRCurrentRetailAveragePrice : 0.0064748932077347485
WarrantyCost : 0.00595886682981698
MMRAcquisitionAuctionAveragePrice : 0.00542409748998343
MMRAcquisitonRetailCleanPrice : 0.004719512839399656
Size : 0.00460513659508568
Nationality : 0.0042393168642713384
MMRAcquisitionAuctionCleanPrice : 0.0030628930478118353
VNST : 0.002931243056199737
Color : 0.002578750625361529
MMRCurrentAuctionAveragePrice : 0.001990652939404491

```

1. What is the significant difference do you see between these two decision tree models (steps 2.1 & 2.2)? How do they compare performance-wise? Explain why those changes may have happened.

Answer: Step 2.1 gave us maximal tree whereas step 2.2 gave us the optimal tree with the gini criterion. Using grid searchCV, we are taking n samples of data. Among them we are taking n-1 samples as training data nth sample as test data with different hyperparameters like criterion,max_depth and minimum number of sample leaves. We will do the same thing by changing the test sample until all n samples becomes the test data. Then it will give us the optimal tree with the best hyperparameters with best accuracy.

1. From the better model, can you identify which cars could potential be “kicks”? Can you provide some descriptive summary of those cars? Based on the feature importance the top 5 features plays important role in determining the kicked cars. WheelType
Auction VehYear Make TopThreeAmericanName